

Introduction to Machine Learning (incomplete)

Aswin

Overview

I made this document as a way to learn ML for my Masters' thesis . Its not an original work, but a compilation of scribed notes of the ML course by Dr Sanjoy Dasgupta(UCSD), [SD] which I audited. Some parts are drawn directly/inspired from Andrew NG's Stanford CS229 course notes, [ANG]. My primary reference was 'The Elements of Statistical Learning' by Trevor Hastie, Robert Tibshirani, Jerome Friedman, [HTF]. So there will be considerable overlap with the aforementioned materials. This note might lack mathematical rigor but I have tried to give proofs and supplementary topics wherever necessary. Please write to me if you find any inaccuracies. I hope this proves at least moderately interesting or useful for you.

Supervised Learning

In a typical scenario, we have an outcome measurement, usually quantitative (such as a stock price) or categorical (such as heart attack/no heart attack), that we wish to predict based on a set of features (such as diet and clinical measurements). We have a training set of data, in which we observe the outcome and feature measurements for a set of objects (such as people). Using this data we build a prediction model, or learner, which will enable us to predict the outcome for new unseen objects. A good learner is one that accurately predicts such an outcome. The examples above describe what is called the supervised learning problem. It is called supervised because of the presence of the outcome variable to guide the learning process.[HTF]

Variable Types and Terminology

The outcome measurement which we wish to predict denoted as ***outputs*** depend on a set of variables denoted as ***inputs***. Classically, the *inputs* are independent variables whereas *outputs* are dependent variables. The term *features* will be used interchangeably with inputs.

The *outputs* which we wish to predict can be qualitative or quantitative (as in blood sugar level). When the *outputs* are qualitative (as in spams or not spams), it is referred as categorical or discrete variables and are typically represented numerically by codes, as in -spam or not spam can be coded as -1 or 1. Depending upon the kind of output variable, the prediction task can be of two types: *regression* when we predict quantitative outputs and *classification* when we predict qualitative outputs.

The input variables/features are denoted by $x^{(i)}$ and the space of all such $x^{(i)}$ is X . The output variable that we are trying to predict is denoted as $y^{(i)}$ and the space of all such $y^{(i)}$ is Y . A pair $(x^{(i)}, y^{(i)})$ is called a training example and the dataset, we will be using to learn - a collection of n training examples $\{(x^{(i)}, y^{(i)}); i = 1, 2, \dots, n\}$ - is called a training set. The superscript (i) in the notation is simply an index into the training set.

Linear Regression

Given a vector of inputs, $x^T = (x_{(1)}, x_{(2)}, \dots, x_{(d)})$. We need to know functions/hypotheses h that can approximate y as a linear function of x :

$$h_{\theta}(x) = \theta_0 + \sum_{i=1}^d x_i \theta_i$$

θ_i s are parameters/weight parametrizing the space of linear functions mapping from X to Y . The term θ_0 is the intercept, also known as the *bias* in machine learning. It is convenient to include θ_0 in the vector of weights θ and add constant variable 1 to the vector x , so that

$$h_{\theta}(x) = \theta^T x$$

For a training set, we have to learn the parameters θ , so that we can predict y . One reasonable method seems to be to make $h(x)$ close to y , for the training examples. To formalize this, we will define a function that measures, for each value of the s , how close the $h(x^{(i)})$'s are to the corresponding $y^{(i)}$'s.

We define Cost/Loss function:

$$J(\theta) = \frac{1}{2} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2,$$

this is the Least Squares cost function. In this approach, we pick the coefficients θ to minimize the cost function J . The LS cost function is quadratic function in weights, θ and hence its minimum always exist, but may not be unique. Given a set of training examples $(x^{(i)}, y^{(i)})$ i.e training set, define a matrix X to be the m -by- n matrix that contains the input values of training examples in its rows:

$$\begin{bmatrix} (x^{(1)})^T \\ \vdots \\ (x^{(m)})^T \end{bmatrix}$$

Let \mathbf{y} be the m -dimensional vector containing the target/output values from the training set:

$$\begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

Since $h_\theta(x^{(i)}) = (x^{(i)})^T \theta$, we can verify that

$$\begin{aligned} X\theta - \mathbf{y} &= \begin{bmatrix} (x^{(1)})^T \theta \\ \vdots \\ (x^{(m)})^T \theta \end{bmatrix} - \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix} \\ &= \begin{bmatrix} h_\theta(x^{(1)}) - y^{(1)} \\ \vdots \\ h_\theta(x^{(m)}) - y^{(m)} \end{bmatrix} \end{aligned}$$

$\frac{1}{2}(X\theta - \mathbf{y})^T(X\theta - \mathbf{y}) = \frac{1}{2} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 = J(\theta)$ This is the cost function. To minimize J , we have to find the derivatives with respect to θ

Some matrix derivative results

$$\nabla_A \text{tr} AB = B^T$$

$$\nabla_A |A| = |A|(A^{-1})^T$$

$$\nabla_{A^T} f(A) = (\nabla_A f(A))^T$$

$$\nabla_A \text{tr} ABA^T C = CAB + C^T AB^T$$

Using the last two results

$$\nabla_{A^T} \text{tr} ABA^T C = B^T A^T C^T + BA^T C$$

The cost function, $J(\theta) = \frac{1}{2}(X\theta - \mathbf{y})^T(X\theta - \mathbf{y})$

$$\begin{aligned} \nabla_\theta J(\theta) &= \nabla_\theta \frac{1}{2}(X\theta - \mathbf{y})^T(X\theta - \mathbf{y}) \\ &= \frac{1}{2} \nabla_\theta (\theta^T X^T X \theta - \theta^T X^T \mathbf{y} - \mathbf{y}^T X \theta + \mathbf{y}^T \mathbf{y}) \\ &= \frac{1}{2} \nabla_\theta \text{tr} (\theta^T X^T X \theta - \theta^T X^T \mathbf{y} - \mathbf{y}^T X \theta + \mathbf{y}^T \mathbf{y}) \\ &= \frac{1}{2} \nabla_\theta (\text{tr} \theta^T X^T X \theta - 2 \text{tr} \mathbf{y}^T X \theta) \\ &= \frac{1}{2} (X^T X \theta + X^T X \theta - 2X^T \mathbf{y}) = X^T X \theta - X^T \mathbf{y} \end{aligned}$$

To minimize J , we set the derivative to zero and obtain: $X^T(X\theta - \mathbf{y}) = 0$. If $X^T X$ is nonsingular then the value of θ that minimizes $J(\theta)$ is given in closed form by the equation,

$$\theta = (X^T X)^{-1} X^T \mathbf{y}$$

Now that we have the parameters, we can predict the output corresponding to the input $x^{(i)}$ as $y^{(i)} = \theta^T x^{(i)}$